

Operating System Project #1 Tutorial

Understanding and Implementing
System Calls

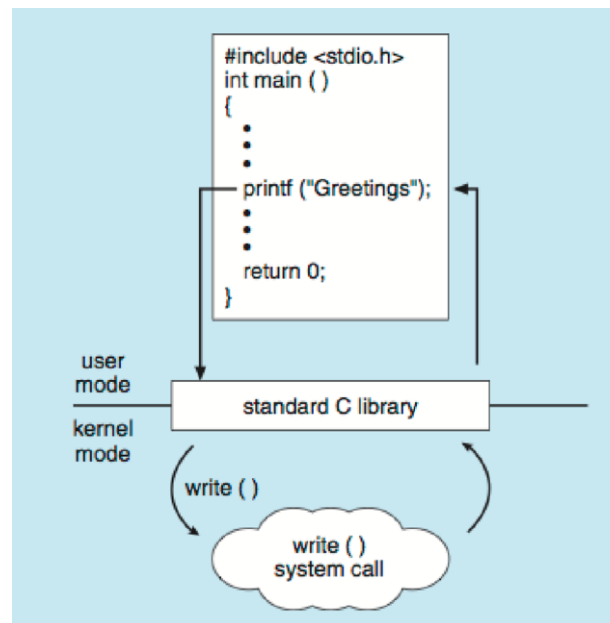
2024. 03. 24

Purpose

- **Modify and compile source code for Linux to add new system calls**
- **Create a user application that uses added system calls**
- **Understand how a system call works on Linux**

System call

- **Path to enter kernel mode from user mode**
 - Required to use the protected (privileged) service provided by the kernel
- **User programs usually use the high-level Application Programming Interface(API) rather than using direct system calls**
 - Provides a more convenient interface for users



Results up to the Project #0

- Install Virtual Box
- Install the Linux on Virtual Machine
- Compile Linux-4.20.11 Kernel

Project #1 is to proceed after the Project #0

Contents of the Project #1

- **Add system call**


- Define a **queue** structure that 1) enqueues integer values and 2) dequeues the integer **in the order it was enqueued**
- **Two system calls** to enqueue and dequeue an integer value

- **Condition**

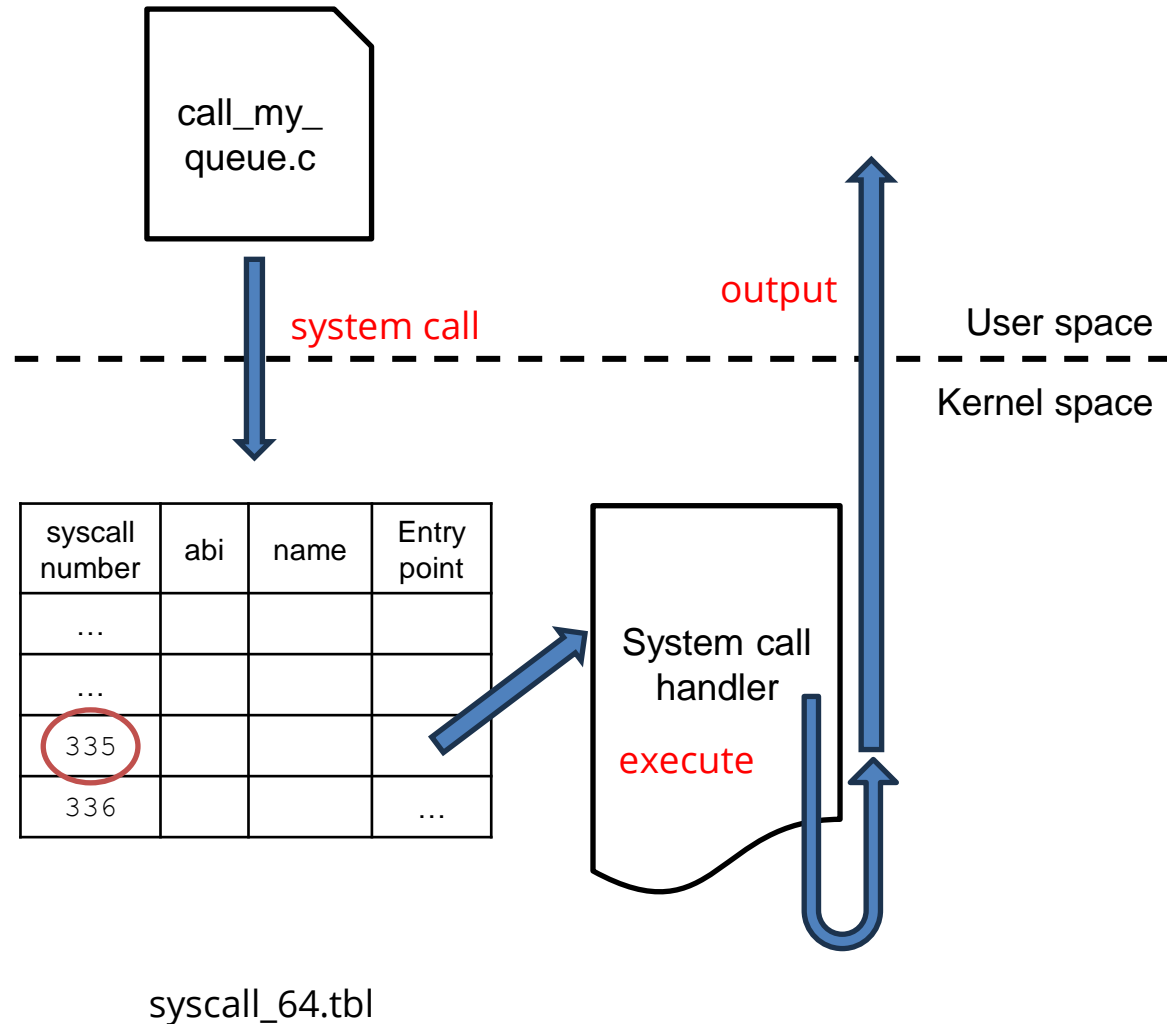
- Enqueue system call takes an “integer” value as its parameter
- When the value to be enqueued already exists in the queue, the system call handler should not newly store the value in the queue
- When dequeue system call is called, its handler pulls the oldest integer off the queue.
 - The dequeued (returned) value is the oldest integer, i.e. the integer enqueued first into the queue.
- User application calls enqueue and dequeue system calls at least three times
 - When calling enqueue system call, it should generate a **random** number and pass it as a parameter

Contents of the Project #1

Compile your kernel



Workflow



Result example

- Application output

```
osta@osta-VirtualBox:~/project1$ ./call_my_queue
Enqueued: 43
Enqueued: 27
Enqueued: 95
Dequeued: 43
Dequeued: 27
Dequeued: 95
```

- Kernel log output (check by dmesg)

```
[ 68.367096] [System call] os2024_enqueue(); -----
[ 68.367096] Queue Front -----
[ 68.367097] 43
[ 68.367097] Queue Rear -----
[ 68.367174] [System call] os2024_enqueue(); -----
[ 68.367174] Queue Front -----
[ 68.367175] 43
[ 68.367175] 27
[ 68.367175] Queue Rear -----
[ 68.367178] [System call] os2024_enqueue(); -----
[ 68.367179] Queue Front -----
[ 68.367179] 43
[ 68.367179] 27
[ 68.367180] 95
[ 68.367180] Queue Rear -----
```

```
[ 68.367182] [System call] os2024_dequeue(); -----
[ 68.367183] Queue Front -----
[ 68.367183] 27
[ 68.367183] 95
[ 68.367184] Queue Rear -----
[ 68.367186] [System call] os2024_dequeue(); -----
[ 68.367187] Queue Front -----
[ 68.367187] 95
[ 68.367187] Queue Rear -----
[ 68.367190] [System call] os2024_dequeue(); -----
[ 68.367190] Queue Front -----
[ 68.367190] Queue Rear -----
```

1. Implementation on Kernel

- **Modify and create the following files specified in the project instruction**

- A. syscall_64.tbl

- A file that collects symbolic information about the names of system call handler (function)
 - **New system call number should be added**

- B. syscalls.h

- Define **prototypes** and register table for added system call functions

- C. sslab_my_queue.c

- Create in /usr/src/linux-4.20.11/kernel/
 - **Source of new system call handlers** to be added

- D. Makefile

- Modifying /usr/src/linux-4.20.11/kernel/Makefile
 - **Add object to be compiled with (sslab_my_queue.o)**

1-A. syscall_64.tbl

- **Maintain unique numbers for all system calls supported by Linux**
 - (linux)/arch/x86/entry/syscalls/syscall_64.tbl
 - ※ (linux) : /usr/src/linux-4.20.11 (Root folder where the kernel's source code is stored)
- **Symbol information table for system calls**
 - Information managed by Linker
 - Table that stores addresses of system call functions located in various spots in Linux kernel code directories
 - Linker automatically manages the system call addresses and locations for kernel compile

```
osta@osta-VirtualBox: /usr/src/linux-4.20.11
File Edit View Search Terminal Help
334      common  rseq                      __x64_sys_rseq
# sslab --- start
335      common  os2024_enqueue                __x64_sys_os2024_enqueue
336      common  os2024_dequeue                 __x64_sys_os2024_dequeue
# sslab --- end

#
# x32-specific system call numbers start at 512 to avoid cache impact
# for native 64-bit operation. The __x32_compat_sys stubs are created
# on-the-fly for compat_sys_*( ) compatibility system calls if X86_X32
```

1-B. syscalls.h

- Define prototypes of system call functions

- Add them to (linux) /include/linux/syscalls.h
- `asmlinkage int sys_os2024_enqueue(int);`
- `asmlinkage int sys_os2024_dequeue(void);`

```
static inline unsigned int ksys_personality(unsigned int personality)
{
    unsigned int old = current->personality;

    if (personality != 0xffffffff)
        set_personality(personality);

    return old;
}

/* sslab --- start */
asmlinkage int sys_os2024_enqueue(int);
asmlinkage int sys_os2024_dequeue(void);
/* sslab --- end */

#endif
#include/linux/syscalls.h" 1304L, 50760C written 1303,0-1
```

- Why use asmlinkage?

- System call calls from trap instructions
- Handling routine for the trap instructions – created with assembly code
- Declaring asmlinkage before function enables calling of C functions within assembly code (trap routines)

1-C. sslab_my_queue.c

- **System call handlers (functions) to add**

- Implement what system calls will actually task (system call handlers)
 - Create in `/usr/src/linux-4.20.11/kernel/`
- Declare a queue in the form of an integer array as **global variable**
- **Skeleton implementation**
 - `SYSCALL_DEFINE1(os2024_enqueue, int, a){`
 `...`
 `}`
 - `SYSCALL_DEFINE0(os2024_dequeue){`
 `...`
 `}`
- `SYSCALL_DEFINEx`: Macros for implementing system calls with “x number of parameters”
 - Defined in `(linux)/include/linux/syscalls.h`

- **Use the following headers**

- `<linux/syscalls.h>`
- `<linux/kernel.h>`
- `<linux/linkage.h>`

1-C. sslab_my_queue.c

```
/* 2024 Spring COSE341 Operating System */
/* Project 1 */
/* your name here */

#include <linux/syscalls.h>
#include <linux/kernel.h>
#include <linux/linkage.h>

SYSCALL_DEFINE1(os2024_enqueue, int, a) {
    /* your implementation here */
}

SYSCALL_DEFINE0(os2024_dequeue) {
    /* your implementation here */
}
```

1-D. Makefile

- /usr/src/linux-4.20.11/kernel/Makefile
- Add things to be compiled within kernel (obj-y part)

```
# SPDX-License-Identifier: GPL-2.0
#
# Makefile for the linux kernel.
#

obj-y      = fork.o exec_domain.o panic.o \
             cpu.o exit.o softirq.o resource.o \
             sysctl.o sysctl_binary.o capability.o ptrace.o user.o \
             signal.o sys.o umh.o workqueue.o pid.o task_work.o \
             extable.o params.o \
             kthread.o sys_ni.o nsproxy.o \
             notifier.o ksysfs.o cred.o reboot.o \
             async.o range.o smpboot.o ucount.o sslab_my_queue.o
```

- .o object file name is the same as its own C file name
 - E.g., sslab_my_queue.c -> sslab_my_queue.o

Kernel Compile

- If you have completed all of the above, run the following commands in the `/usr/src/linux-4.20.11` (kernel source code folder)

```
$ sudo make
```

```
$ sudo make install
```

2. Implmenetation of user application

- Create an application that uses the added system calls

- Use a macro function called **syscall()** to call the system
 - How to use it: pass the system call number and parameters
 - `syscall(335, ...);`
 - » Add header `<unistd.h>`
 - `#define my_queue_enqueue 335` // After declaring the system call number
 - `syscall(my_queue_enqueue, ...);` more convenient to use as

- Application compile

- If the application source file is saved as **call_my_queue.c**
 - `gcc call_my_queue.c -o call_my_queue`
 - Means that “compile `call_my_queue.c` and create an executable (e.g., `call_my_queue`)”
 - Run to the `./call_my_queue` and check the `dmesg` to see if the desired output has come from by `printk` as well

Note that code changes may not apply if they are not Linux 4.20.11 kernels.
If the system call you added is not loaded, check the kernel version by the command `'uname -r'`.
Then, reboot, press “left shift key”, and select the proper kernel version (the one you implemented your system calls) in Advanced options for Ubuntu

Project submission

- **List of files to submit (Check the Project #1_instruction.pdf)**
 - Report
 - See **Project #1_instruction.pdf** for the details to explain in the report
 - All source code files you modified or newly implemented
 - Comments explaining the roles of important variables and functions should be included in the source code
 - Execution result file
 - Make the execution results into the text log (result.txt) and submit it

Tip

- **Most questions can be solved through web searching**

- There are a lot of web documents related to the Linux kernel
- To use your time ***effectively***, first, look for solutions online.
If you encounter specific issues you can't resolve after searching, do not hesitate to ask questions by contacting the TAs.

- **Individual Q&A**

- Blackboard Q&A Bulletin
- Email: osta@os.korea.ac.kr